Quiz 6

Problem 1

1/1 point (graded)

You are given a binary 4-dimensional linear decision boundary with coefficient vector $\mathbf{w} = [2, 1, 4, 3]$ and $b = -12$. How would you classify the point $(2, 1, 1, 2)$?

- [ ] -1
- [ ] 0
- [x] 1

✔

Submit

Problem 2

1/1 point (graded)

In which of the following situations has our linear classifier correctly labeled a data point? Select all that apply.

- [x] $\mathbf{w} \cdot \mathbf{x} + b > 0$ and $y > 0$
- [x] $y(\mathbf{w} \cdot \mathbf{x} + b) > 0$
- [x] $\mathbf{w} \cdot \mathbf{x} + b < 0$ and $y < 0$
- [ ] $y > \mathbf{w} \cdot \mathbf{x} + b$

Generating Speech Output

Submit

Problem 3

1/1 point (graded)

Let's say that we have a linear classifier given by $\mathbf{w} = [1, 1, -3, 0]$ and $b = -2$. Our loss function measures the amount by which our prediction is incorrect: $\text{loss} = -y\left(\mathbf{w} \cdot \mathbf{x} + b\right)$. If our prediction is correct, there is no loss.

What is the loss on the data point $(\mathbf{x}, y)$ where $\mathbf{x} = (3, 1, 1, 4)$ and $y = 1$?

- 0
- 1 ✓
- 2
- 3

Submit

Problem 4

1/1 point (graded)

If the Perceptron algorithm does 9 updates before converging on a solution, what value of b can you expect to have?

- $b = 9$
- $b = -9$
- $b \in [-9, 9]$ ✓

Generating Speech Output

- $b \in [0, 9]$

Submit

Problem 5

1/1 point (graded)
What is a support vector?

- A data point from the test set that is used to test the classifier
- A vector that we are trying to minimize
- A data point which is correctly classified by the optimal solution for $\mathbf{w}$
- A data point from the training set that contributes to the optimal solution for $\mathbf{w}$

Submit

Problem 6

1/1 point (graded)
What is the difference between the perceptron algorithm and the support vector machine?

- The perceptron uses gradient descent while the SVM uses stochastic gradient descent
- The perceptron finds a linear separator that separates most of the data points in the training set, while a SVM finds a linear separator that separates all of the data in the training set

Generating Speech Output

- (●) The perceptron finds any solution that perfectly separates the training set, while the SVM finds the solution that perfectly separates the training set with the greatest margin of separation
- ( ) The perceptron algorithm may not find a solution while the SVM is guaranteed to find a solution

✔

Submit

Problem 7

1/1 point (graded)

The optimal solution for a SVM is given by the coefficient vector $\mathbf{w}$ and the constant b. The width of the margin is given by γ. What is the value of γ?

- (●) $\gamma = \frac{1}{||\mathbf{w}||}$
- ( ) $\gamma = \frac{1}{||\mathbf{w}||}$
- ( ) $\gamma = ||\mathbf{w}||$
- ( ) $\gamma = b - \frac{1}{||\mathbf{w}||}$

✔

Submit

Problem 8

1/1 point (graded)

True or false: A soft-margin SVM has fewer support vectors than a hard-margin SVM.

- ( ) True

Generating Speech Output

- (•) False

✔

Submit

Problem 9

1/1 point (graded)

Decreasing the value of C in the soft-margin SVM results in which of the following:

- [ ] fewer number of support vectors
- [x] wider margin
- [ ] more data points being correctly classified
- [x] lower penalty for incorrectly classified data points

✔

Submit

Problem 10

1/1 point (graded)

True or false: All support vectors are contained between, or on, the margins of the two classes.

- ( ) True
- (•) False

✔

Submit

Generating Speech Output

Problem 11

1/1 point (graded)
What does the slack variable represent?

- (●) It is a vector containing the amount of error each point $(x^{(i)}, y^{(i)})$ contributes to the optimization problem
- ( ) It is a coefficient that we must determine to optimize the problem
- ( ) It is a vector containing the number of times each w_i is updated
- ( ) it is a value that determines how much error the optimization problem is allowed to have

✔

Submit

Problem 12

1/1 point (graded)
Using the dual form of the perceptron algorithm, which of the following values are updated during each pass over the training set?

- [ ] $\mathbf{w}$
- [x] α
- [x] b
- [ ] $\mathbf{y}$

✔

Submit

Generating Speech Output

Problem 13

1/1 point (graded)

When optimizing the dual form of the hard-margin SVM, when are the values α_i non-zero?

- When the data point $(x^{(i)}, y^{(i)})$ is on the linear separator between the two classes
- **When the data point $(x^{(i)}, y^{(i)})$ is right on the margin for its class** ✔
- When the data point $(x^{(i)}, y^{(i)})$ is in the interior of the region for its class
- When the data point $(x^{(i)}, y^{(i)})$ is on the wrong side of the linear separator

Submit

Problem 14

1/1 point (graded)

When using multiclass logistic regression on data with labels, $Y = \{1, 2, \ldots, k\}$, and a linear classifier specified by $\mathbf{w}_1, \mathbf{w}_2, \ldots, \mathbf{w}_k \in \mathbb{R}^d$ and $b_1, b_2, \ldots, b_k \in \mathbb{R}$, and given a point $(\mathbf{x}, y)$, what is the probability that $y = j$, where $0 < j \leq k$?

- $Pr(y = j|\mathbf{x}) = e^{\mathbf{w}_j \cdot \mathbf{x} + b_j}$
- $Pr(y = j|\mathbf{x}) = \frac{e^{\mathbf{w}_j \cdot \mathbf{x} + b_j}}{e^{\mathbf{w}_k \cdot \mathbf{x} + b_k}}$
- **$Pr(y = j|\mathbf{x}) = \frac{e^{\mathbf{w}_j \cdot \mathbf{x} + b_j}}{e^{\mathbf{w}_1 \cdot \mathbf{x} + b_1} + e^{\mathbf{w}_2 \cdot \mathbf{x} + b_2} + \ldots + e^{\mathbf{w}_k \cdot \mathbf{x} + b_k}}$** ✔
- $Pr(y = j|\mathbf{x}) = \frac{e^{\mathbf{w}_j \cdot \mathbf{x} + b_j}}{1 + e^{\mathbf{w}_j \cdot \mathbf{x} + b_j}}$

Generating Speech Output

Submit

Problem 15

1/1 point (graded)

What does ξ_i represent in the soft-margin SVM?

- It is the number of times the i'th point was updated
- It is the amount of slack the i'th point has
- It represents the i'th support vector
- It represents the width of the margin

Submit

Generating Speech Output