



[Course](#) > [Week 4...](#) > [Progra...](#) > Progra...

Programming Assignment 4

Click this link to download the [Diabetes Regression notebook](#) and then complete problems 1-4.

Click this link to download the [mystery.dat file](#) which will help you complete problem 5.

Click this link to download the [Sentiment Logistic Regression notebook](#).

Problem 1

1/1 point (graded)

This problem is based on the *Diabetes Regression notebook*. You should work through that notebook before entering your answers here.

If a single feature is to be used to predict y , the best choice (the one that yields the smallest MSE) is feature 2 ('body mass index'). What is the second-best choice? Your answer should be the feature number (0 – 9).



Submit

Problem 2

2/2 points (graded)

Use the _____ procedure to create training/test splits of various sizes. In particular, try training set sizes of 20, 50, 100, and 200. In each case, record the training error and test error *when using all features for prediction*.

For a training set size of 100, what are the training MSE and test MSE (just round to the nearest integer)?

Training MSE =

2883.7785



Test MSE =

3583.00



Submit

Problem 3

1/1 point (graded)

What *rough* trends do you observe as the training set size increases (from, say, 20 to 400)? Select all that apply.

☒ The training error increases

☒ The test error decreases

☒ The gap between the training and test error decreases



Submit

Problem 4

1/1 point (graded)

What is the single best explanation for these trends? Choose one of the following.

☐ With more training data, we get better estimates of training error.

☒ With more training data, we learn a more accurate model.

☐ The error is proportional to the amount of data.



Submit

Problem 5 relates to finding relevant features.

Problem 5

1/1 point (graded)

The file contains pairs (x, y) , where $x \in \mathbb{R}^{100}$ and $y \in \mathbb{R}$. There is one data point per line, with comma-separated values; the very last number in each line is the y -value.

In this data set, y is a linear function of just *ten* of the features in x , plus some noise. Your job is to identify those ten features.

Which of the following contain only relevant features?

(Think of the feature numbers as being in the range 1 to 100, but be aware that Python indexes arrays starting at zero.)

☐ 1,5,7,19,44

☒ 2,3,13,17,29

☐ 3,7,13,19,44

☐ 5,23,24,51,61



Submit

